

# Lab 3 - Validating Linearizability of Lock-free Skiplists

- Group 18
- Wenqi Cao

## 1. Measuring execution time

### 1.1 Measurement program

I modified the measurement program as follows for testing.


```
java Main 2 Default Normal 4096 1:1:8 100000 20 50
```

Operations with 50% add & 50% remove seem to be slower.

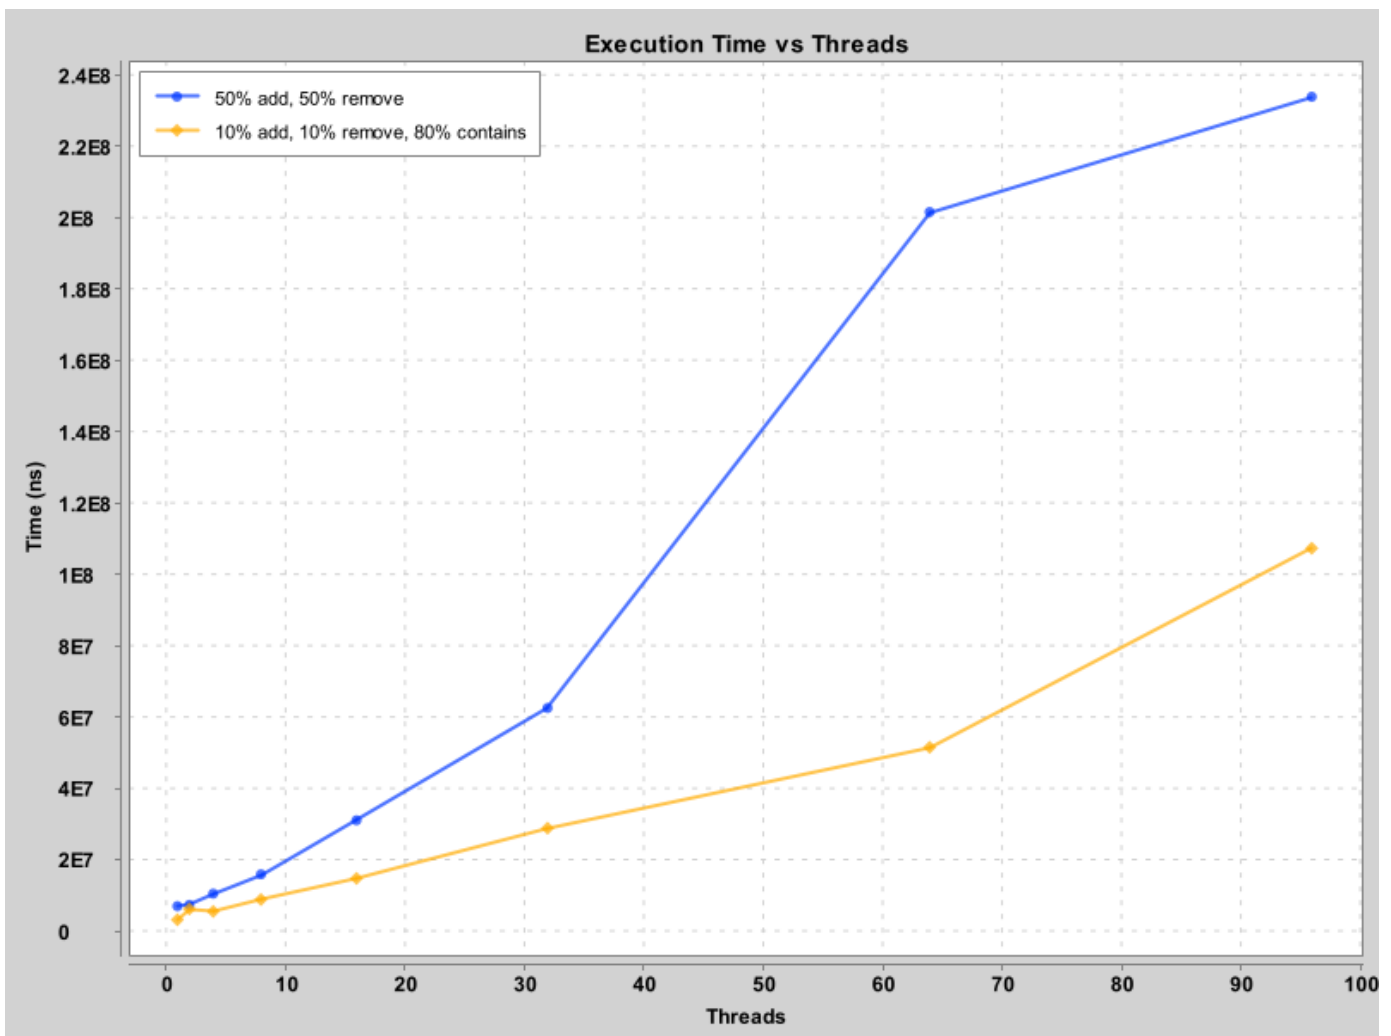
### 1.2 Dardel experiments

Source file:

- `src/Main.java` (Run on PDC)
  - `plots/PDCPlot1.java` (Run on PDC)
-



- Fig above is the Normal Distribution.
- Fig below is the Uniform Distribution.



Yes, they make sense.

Uniform distribution is expected to be faster. In normal distribution, the concentration of operations in a small range of values causes contention being high and slows down performance.

The 10:10:80 distribution should be faster than the 50:50 distribution because read-heavy workloads tend to have better performance in concurrent lock-free data structures like skip lists, as they involve less contention and fewer expensive operations.

## 2. Identify and validate linearization points

### 2.1 Identify linearization points

Generally speaking, the locations of these points are around the successful or unsuccessful call.

- `add()` : The linearization point is where the node is successfully inserted into the list with `compareAndSet()` or when it is determined that the node already exists.
- `remove()` : The linearization point is where the node is marked logically deleted or when it is found that the node is already removed or doesn't exist.
- `contains()` : The linearization point is when the element is found in the list or determined to not be present.

Well, if described in my own words, I would say that capture it "before the return".

## 2.2 Develop a validation method

Source file:

- `src/Main.java`
- `src/log.java`

`Log.validate` is implemented with the help of `HashSet` .

## 2.3. Locked time sampling

Source file:

- `src/Main.java`
- `src/LockFreeSkipListLocked.java`

Though the locked version is more accurate, it introduces large delays, especially as the number of threads increases. The lock contention causes performance to degrade in multithreading scenarios.

## 2.4. Lock-free time sampling with local log

Source file:

- `src/Main.java`
- `src/LockFreeSkipListLocalLog.java`

## 2.5. Lock-free Time Sampling with Global Log

Source file:

- `src/Main.java`
- `src/LockFreeSkipListGlobalLog.java`

The absence of locks means reduced contention and improved throughput. However, it brings some trade-offs in accuracy, particularly in the ordering of timestamps due to the lack of precise synchronization between threads.

## 2.5.Extra

Source file:

- `src/LockFreeQueue.java`


Reference: HSLs Chapter 10 Page 237-238 LockFreeQueue.

## 2.6. Dardel experiments


Source file:

- `plots/PDCPlot1.java` (Run on PDC)
  - `plots/PDCPlot2.java` (Run on PDC)
-


Execution Time vs Threads



Discrepancy vs Threads




Execution Time vs Threads







Discrepancy vs Threads




Execution Time vs Threads



Discrepancy vs Threads



Execution Time vs Threads



Discrepancy vs Threads

